

Machine Learning:

Logistic Regression

Xiuxia Du, Ph.D.
Department of Bioinformatics and Genomics
University of North Carolina at Charlotte

Outline

- Introduction
- Odds and Odds Ratio
- Logistic Regression Core Idea
- Build Logistic Regression Models
- Interpret coefficients
- How to learn the coefficients?
- Regularization

Introduction

Classification

- Make a decision:
 - Email: Spam / Not Spam?
 - Tumor: Malignant / Benign?
- Mathematical representation:
 - $y \in 0, 1$ with
 - $y = 0$: Negative class (e.g, benign tumor)
 - $y = 1$: Positive class (e.g., malignant tumor)

Odds and Odds Ratio

Probability

- Probability (p)
 - $p = P(Y = 1)$
 - It always lies between 0 and 1.
- Example
 - If $p = 0.8$, there is an 80% chance of success.

Odds

- Odds are defined as:

$$\text{odds} = \frac{p}{1 - p}$$

- So odds compare “how likely somethings is to happen” vs. “how likely it is NOT to happen”

Example:

Probability p	Odds $\frac{p}{1 - p}$	Interpretation
0.9	9	“9 to 1” in favor
0.8	4	“4 to 1” in favor
0.5	1	equally likely, yes vs no
0.2	0.25	“1 to 4” in favor

Odds

Key intuition:

- Odds = 1 means a 50/50 event.
- Odds > 1 means event is more likely than not.
- Odds < 1 means event is less likely than not.

Odds Ratio (OR)

- Definition: An odds ratio compared the odds between two groups:

$$\text{odds ratio} = \frac{\text{odds in group A}}{\text{odds in group B}}$$

- So it answers the question:

How many times larger are the odds in group A than in group B?

Odds Ratio

Example

- Suppose:
 - In the control group, $p = 0.2$.
 - In the treatment group, $p = 0.5$.

- Compute odds

- Control group:

$$\text{odds} = \frac{0.2}{0.8} = 0.25$$

- Treatment group:

$$\text{odds} = \frac{0.5}{0.5} = 1$$

- Compute odds ratio:

$$\text{odds ratio} = \frac{1}{0.25} = 4$$

- **Interpretation:** The treatment group has 4x higher odds of success than the control group.

Odds Ratio

Odds Ratio is NOT the same as probability ratio.

- Probability ratio (risk ratio)

$$\frac{p_1}{p_0}$$

where $p_1 = p(Y = 1|x; \theta)$ for group 1 and $p_0 = p(Y = 1|x; \theta)$ for group 0.

- Odds ratio between group 1 and group 2

$$\frac{p_1/(1 - p_1)}{p_0/(1 - p_0)}$$

- These are only similar when probabilities are small.

Odds Ratio: Example

Odds Ratio is NOT the same as probability ratio.

- Probability

- Control: $p = 0.5$
- Treatment: $p = 0.75$

- Risk ratio:

$$\frac{0.75}{0.5} = 1.5$$

- Odds

- Control odds = $0.5/0.5 = 1$
- Treatment odds = $0.75/0.25 = 3$

- Odds ratio

$$\frac{3}{1} = 3$$

There is a big difference between the risk ratio and odds ratio.

Logistic Regression Core Idea

Core Idea

Procedure


- Linearly combine the independent variables

$$z = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_d x_d$$

where $z \in (-\infty, \infty)$.

- Sigmoid function to convert the real number z to $g(z) \in (0, 1)$:

$$g(z) = \frac{1}{1 + e^{-z}}$$



Core Idea

Procedure

- The full model:

$$P(y = 1|x) = p = g(z) = \frac{1}{1 + e^{-z}}$$

- This means:

$$z = \ln \frac{p}{1 - p}$$

Core Idea

- Odds

$$\text{odds} = \frac{p}{1-p}$$

- log-odds (logit)

$$\ln \left(\frac{p}{1-p} \right)$$

Logistic regression is linear in log-odds:

$$\ln \left(\frac{p}{1-p} \right) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_d x_d$$

Build Logistic Regression Models

Logistic regression model

- We want


$$0 \leq h_{\theta}(x) \leq 1$$

- This can be achieved by using the logistic/sigmoid function:

$$h_{\theta}(x) = g(\theta_0 x_1 + \theta_1 x_1 + \cdots + \theta_n x_n) = g(\theta^T x)$$

with $g(z)$ defined as

$$g(z) = \frac{1}{1 + e^{-z}}$$



Interpretation of logistic regression output

- $h_{\theta}(x)$ is interpreted as the estimated probability that $y = 1$ given x , i.e.

$$\begin{aligned}h_{\theta}(x) &= g(\theta^T x) = P(y = 1|x; \theta) \\1 - h_{\theta}(x) &= g(\theta^T x) = P(y = 0|x; \theta)\end{aligned}$$

- **Example:** If $z = \theta_0 + \theta_1 x$ where x is the tumor size and $h_{\theta}(x) = 0.7$, then we can tell the patient that there is a 70% probability that the tumor is malignant.
- Suppose we predict $y = 1$ when $h_{\theta}(x) \geq 0.5$ and predict $y = 0$ when $h_{\theta}(x) < 0.5$. This means that

$$y = \begin{cases} 1 & \text{if } \theta^T x \geq 0 \\ 0 & \text{if } \theta^T x < 0 \end{cases}$$


Decision boundary

- $\theta^T x = 0$ divides the feature space into two parts. The part where $\theta^T x \geq 0$ corresponds to $y = 1$ and the part where $\theta^T x < 0$ corresponds to $y = 0$.
- **Example:** Suppose $h_\theta(x) = g(\theta_0 + \theta_1 x + \theta_2 x)$ with $\theta_0 = -3, \theta_1 = 1, \theta_2 = 1$. Then


$\theta^T x = 0$ corresponds to the magenta line

$\theta^T x > 0$ corresponds to the area above the magenta line


$\theta^T x < 0$ corresponds to the area below the magenta line



Decision boundary



feature space



z space

Decision boundary

- As x changes, z changes linearly.
- Sigmoid maps z on the real line $(-\infty, +\infty)$ into probability $(0, 1)$.
- Decision boundary occurs when $z = 0$ and $p = 0.5$, i.e.

$$\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_d x_d = 0$$

- Log-odds converts probability $(0, 1)$ to the real line $(-\infty, +\infty)$

$$\ln \left(\frac{p}{1-p} \right) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_d x_d$$


Nonlinear decision boundary

Example: Suppose

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

where $\theta_0 = -1, \theta_1 = \theta_2 = 0, \theta_3 = \theta_4 = 1$. Then


- **Predict $y = 1$ if $x_1^2 + x_2^2 \geq 1$**
- **Predict $y = 0$ if $x_1^2 + x_2^2 < 1$**



Complex nonlinear decision boundary

- More complex nonlinear decision boundaries can be achieved by including higher order terms:

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_1^2 x_2 + \theta_5 x_1^2 x_2^2 + \theta_6 x_1^3 x_2 + \dots)$$



Interpreting Coefficients

Interpreting Coefficients

(1) Logistic regression equation

$$\ln \left(\frac{p}{1-p} \right) = \theta_0 + \theta_1 x_1$$

with

- $p = P(y = 1|x)$
- $\frac{p}{1-p}$ is the odds.

(2) Exponentiate both sides

$$\frac{p}{1-p} = e^{\theta_0 + \theta_1 x_1}$$

So the odds are:

$$\text{odds}(x) = e^{\theta_0 + \theta_1 x_1}$$

Interpreting Coefficients

(3) Compare odds at two x -values

Odds at x :

$$\text{odds}(x) = e^{\theta_0 + \theta_1 x}$$

Odds at $x + 1$:

$$\text{odds}(x + 1) = e^{\theta_0 + \theta_1 (x+1)} = e^{\theta_0 + \theta_1 x + \theta_1}$$

(4) Take the ratio

$$\frac{\text{odds}(x + 1)}{\text{odds}(x)} = \frac{e^{\theta_0 + \theta_1 x + \theta_1}}{e^{\theta_0 + \theta_1 x}} = e^{\theta_1}$$

(5) Final interpretation:

$$e^{\theta_1} = \text{odds ratio for a 1-unit increase in } x$$


Interpreting Coefficients

Interpretation of θ_j :

- Each θ_j affects the log-odds.
- $\theta_j > 0$: increases log-odds $\rightarrow$ higher probability
- $\theta_j < 0$: decreases log-odds $\rightarrow$ lower probability

Or we can say:

- If $e^{\theta_1} = 1$: no effect
- If $e^{\theta_1} > 1$: increases odds
- If $e^{\theta_1} < 1$: decreases odds



Example

Given a 1-feature logistic regression model:

$$\ln \left(\frac{p}{1-p} \right) = \theta_0 + \theta_1 x$$

where $\theta_0 = 0$ and $\theta_1 = \ln(2)$.

So the model is:

$$\ln \left(\frac{p}{1-p} \right) = \ln(2) \cdot x$$

Interpreting Coefficients: Example

Step 1: Compute the odds when $x = 0$

- Log-odds:

$$\ln \left(\frac{p}{1-p} \right) = \ln(2) \cdot 0$$

- Odds:

$$\frac{p}{1-p} = e^0 = 1$$

- So:

$$\text{odds}(0) = 1$$

This corresponds to $p = 0.5$.

Interpreting Coefficients: Example

Step 2: Compute the odds when $x = 1$

- Log-odds:

$$\ln \left(\frac{p}{1-p} \right) = \ln(2) \cdot 1 = \ln 2$$

- Odds:

$$\frac{p}{1-p} = e^{\ln(2)} = 2$$

- So:

$$\text{odds}(1) = 2$$

This corresponds to $p = \frac{2}{1+2} = \frac{2}{3} \approx 0.667$.

Interpreting Coefficients: Example

Step 3: Compute the odds ratio

$$\text{odds ratio} = \frac{\text{odds}(1)}{\text{odds}(0)} = \frac{2}{1} = 2$$

Step 4: Compare with e^{θ_1}

$$e^{\theta_1} = e^{\ln(2)} = 2$$

Final interpretation:

When $\theta_1 = \ln(2)$, every time x increases by 1 unit, the odds ratio is 2 and the odds multiply by 2.

Interpreting Coefficients: Example

- Now let's add when $x = 2$ and $x = 3$ for the exact same model with $\theta_0 = 0$ and $\theta_1 = \ln(2)$. So

$$e^{\theta_1} = e^{\ln(2)} = 2$$

Meaning that each +1 increase in x multiplies the odds by 2.

- Compute odds for $x = 0, 1, 2, 3$.

$$\text{odds}(x) = e^{\ln(2)x} = 2^x$$

So:

x	log-odds	odds
0	0	$2^0 = 1$
1	$\ln(2)$	$2^1 = 2$
2	$2 \ln(2)$	$2^2 = 4$
3	$3 \ln(2)$	$2^3 = 8$

Interpreting Coefficients: Example

- Convert odds to probability

$$p = \frac{\text{odds}}{1 + \text{odds}}$$

x	log-odds	odds	probability p
0	0	$2^0 = 1$	$\frac{1}{2} = 0.5$
1	$\ln(2)$	$2^1 = 2$	$\frac{2}{3} \approx 0.667$
2	$2 \ln(2)$	$2^2 = 4$	$\frac{4}{5} = 0.8$
3	$3 \ln(2)$	$2^3 = 8$	$\frac{8}{9} \approx 0.889$

Interpreting Coefficients: Example

- Now compute the odds ratio between consecutive x values:

$$\frac{\text{odds}(x + 1)}{\text{odds}(x)}$$

From $x \rightarrow x + 1$	odds ratio
$0 \rightarrow 1$	$\frac{2}{1} = 2$
$1 \rightarrow 2$	$\frac{4}{2} = 2$
$2 \rightarrow 3$	$\frac{8}{4} = 2$

So the odds ratio is always 2, which equals:

$$e^{\beta_1} = e^{\ln(2)} = 2$$

Interpreting Coefficients: Example

Key points

- Odds increase by a constant factor ($\times 2$)
 - $1 \rightarrow 2 \rightarrow 4 \rightarrow 8$
 - multiplies by 2 each step
- But probabilities do NOT increase by a constant amount.
 - $0.5 \rightarrow 0.667 \rightarrow 0.8 \rightarrow 0.889$
 - increases shrink as you approach 1.
- This is why logistic regression behaves realistically:
 - it has diminishing returns near probability 0 or 1
 - but a consistent effect in log-odds space.

Interpreting Coefficients: Example

Now let's take a look at an example where θ_1 is negative.

- We use the same logistic regression form:

$$\ln \left(\frac{p}{1-p} \right) = \theta_0 + \theta_1 x$$

with $\theta_0 = 0$ and $\theta_1 = \ln(0.5) = -\ln(2)$. So θ_1 is negative.

- Compute the odds ratio

$$e^{\theta_1} = e^{\ln(0.5)} = 0.5$$

meaning that each +1 increase in x multiplies the odds by 0.5, i.e. cuts the odds in half.

Interpreting Coefficients: Example

- Compute odds

$$\text{odds}(x) = e^{\ln(0.5) \cdot x} = (0.5)^x$$

x	log-odds	odds	probability p
0	0	$(0.5)^0 = 1$	$\frac{1}{2} = 0.5$
1	$\ln(0.5)$	$(0.5)^1 = 0.5$	$\frac{0.5}{1.5} \approx 0.333$
2	$2 \ln(0.5)$	$(0.5)^2 = 0.25$	$\frac{0.25}{1.25} = 0.2$
3	$3 \ln(0.5)$	$(0.5)^3 = 0.125$	$\frac{0.125}{1.125} \approx 0.111$

Interpreting Coefficients: Example

- Compute odds ratio

$$\text{odds ratio} = \frac{\text{odds}(x+1)}{\text{odds}(x)}$$

From $x \rightarrow x+1$	odds ratio
$0 \rightarrow 1$	$\frac{0.5}{1} = 0.5$
$1 \rightarrow 2$	$\frac{0.25}{0.5} = 0.5$
$2 \rightarrow 3$	$\frac{0.125}{0.25} = 0.5$

So the odds ratio is always 0.5, which equals:

$$e^{\beta_1} = e^{-\ln(2)} = 0.5$$

meaning that each 1-unit increase in x cuts the odds of class 1 in half.

Interpreting Coefficients

When predictors have been standardized with `StandardScaler` in `sklearn`, the interpretation of the coefficients is:

Holding all other features constant, a one-standard-deviation increase in feature X_j is associated with a multiplicative change of e^{θ_j} in the odds of the outcome being class 1.

How to Learn the Coefficients?

Cost function

- The training set consists of m examples:

$$\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$$

- Let

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}$$

with $x_0 = 1, y \in 0, 1$.




- Use the following logistic function to determine y :

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Cost function for linear regression

- Cost function for linear regression $h_{\theta}(x) = \theta_0 + \theta_1 x$:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$



- However, this type of cost function is not appropriate for logistic regression since the cost function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (g(\theta_0 + \theta_1 x^{(i)}) - y^{(i)})^2$$

with $g(\cdot)$ being the logistic function might not be convex.



Cost function for logistic regression

- Define

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

with

$$\text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) = \begin{cases} -\log(h_{\theta}(x^i)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x^{(i)})) & \text{if } y = 0 \end{cases}$$



Cost function for logistic regression

- The logistic cost function defined for $y = 0$ and $y = 1$ can be combined into one equation as follows:

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

- To fit parameters θ , we minimize $J(\theta)$ to get

$$\text{argmin}_{\theta} J(\theta)$$

- After training, predictions can be made for a new given x as:

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

that is interpreted as $P(y = 1|x; \theta)$.

Gradient descent

- To minimize

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right]$$

We simultaneously update

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

and repeat the procedure until convergence. The partial derivative term is:

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)}$$

Gradient descent for linear regression

- Linear model:

$$h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \cdots + \theta_n x_n$$

- Cost function:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

- Gradient descent by simultaneously updating θ and repeat the procedure until convergence:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) = \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)}$$

Gradient descent for linear regression

- More specifically:

$$\theta_0 \quad := \quad \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta) = \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x_0^{(i)}$$

$$\theta_1 \quad := \quad \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta) = \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x_1^{(i)}$$

$\vdots$

$$\theta_n \quad := \quad \theta_n - \alpha \frac{\partial}{\partial \theta_n} J(\theta) = \theta_n - \alpha \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right) x_n^{(i)}$$

Regularization

What is regularization?

- Regularization discourages a model from using excessively large coefficients, making the model less complex and often improving its ability to generalize to new data.
- Consider logistic regression:

$$\log \left(\frac{p}{1-p} \right) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_d x_d$$

Without regularization, logistic regression chooses the coefficients $\theta_0, \theta_1, \dots, \theta_d$ that best fit the training data. The problem is that the model can sometimes fit the training data **too well**.

Why can large coefficients be a problem?

Imagine predicting disease status using three biomarkers:

$$\text{logit}(p) = -1.2 + 0.2x_1 - 0.8x_2 + 18.7x_3$$

The coefficient for x_3 is extremely large. That may mean x_3 is genuinely an extremely strong predictor. But it could also mean that the model has found a peculiar pattern in the training data. For example, perhaps by chance the following happens in the training data:

$x_3 < 1.5 \rightarrow$ almost everyone healthy

$x_3 > 1.5 \rightarrow$ almost everyone diseases

The model may assign x_3 a very large coefficient to exploit this pattern. But the pattern might not hold in new patients. This is **overfitting**.

What does regularization do?

- Regularization modifies the model's objective.
- Without regularization, conceptually the model minimizes: Loss
We ask: How well does the model classify the training observations?
- With regularization, it minimizes: $\text{Loss} + \lambda \times \text{Penalty}$
We ask: How large or complex are the coefficients?
- So the optimizer must make a compromise. A set of coefficients might fit the training data slightly better but require enormous coefficient values. Regularization says:
I will accept a slightly worse training fit if I can obtain a simpler model with smaller coefficients.
That tradeoff is the heart of regularization.

L2 regularization

- A common form of regularization for logistic regression is **L2 regularization**.
- The penalty is based on squared coefficients:

$$\theta_1^2 + \theta_2^2 + \cdots + \theta_d^2$$

So the optimization becomes conceptually

$$\text{Loss} + \lambda \sum_{j=1}^d \theta_j^2$$

- Suppose: $\theta_1 = 1, \theta_2 = 2, \theta_3 = 10$
The L2 penalty is: $1^2 + 2^2 + 10^2 = 105$
That coefficient of 10 contributes: $10^2 = 100$ by itself.
- Therefore, L2 strongly discourages very large coefficients.

What happens to the coefficients?

- Imagine the unregularized model produces:

$$[0.8, -1.3, 7.5, 0.2]$$

With increasing L2 regularization, we might see:

$$[0.7, -1.1, 4.2, 0.18]$$

then

$$[0.4, -0.7, 1.8, 0.10]$$

- The coefficients are shrunk toward zero. This is why regularization is also called **shrinkage**.
- Importantly, L2 generally does not make coefficients exactly zero.

L1 regularization

- L1 regularization uses the absolute values of coefficients:

$$|\theta_1| + |\theta_2| + \cdots + |\theta_d|$$

Conceptually: $\text{Loss} + \lambda \sum_{j=1}^d |\theta_j|$

- A major difference between L1 and L2 regularization is that L1 can drive coefficients exactly to zero.

- For example:

No regularization: $[0.8, -1.3, 0.15, 0.07, 2.1]$

L1 regularization: $[0.5, -0.9, 0, 0, 1.4]$

Now variables 3 and 4 have effectively disappeared from the model.

- Therefore, **L1 regularization can perform implicit feature selection**. This can be attractive in high-dimensional biological datasets such as gene expression, metabolomics, proteomics, or exposomics.

Elastic Net combines L1 and L2

- Elastic Net combines the two:

$$\lambda \left[\alpha \sum_j |\theta_j| + (1 - \alpha) \sum_j \theta_j^2 \right]$$

In current scikit-learn terminology, `l1_ratio` plays the role of α :

$$\text{l1_ratio} = 0 \Rightarrow L2$$

$$\text{l1_ratio} = 1 \Rightarrow L1$$

$$\text{l1_ratio} = 0.5 \Rightarrow \text{mixes L1 and L2}$$

- So Elastic Net can simultaneously shrink coefficients and eliminate some variables.

C

- Many textbooks write regularization strength as C .
- Scikit-learn parameterizes logistic regression using C , which is inverse regularization strength, i.e.

$$C \propto \frac{1}{\lambda}$$

- Therefore

small $C \rightarrow$ strong regularization $\rightarrow$ More coefficient shrinkage $\rightarrow$ simpler model

whereas:

large $C \rightarrow$ weak regularization $\rightarrow$ less coefficient shrinkage $\rightarrow$ more flexible model

- In scikit-learn, use `LogisticRegression(C=100)` or `LogisticRegression(C=0.001)`. The latter uses a much stronger regularization.

Bias-variance tradeoff

- Too little regularization: The model becomes flexible:

low bias + high variance

It may fit peculiarities of the training data and cause **overfitting**.

- Too much regularization: The coefficients become excessively constrained and

high bias + low variance

The model cannot capture genuine relationships and causes **underfitting**.

- We want a balance and appropriate regularization. So we need to tune C .



Why scaling is important for regularization

- Suppose two predictors are:
 - Age: 20 - 80
 - Protein level: 0 - 100,000
- The numerical scale affects the size of the coefficients needed to produce a given change in the linear predictor.
- Because regularization penalizes **coefficient magnitude**, predictors measured on different scales can be penalized unfairly.
- Therefore, when using regularized logistic regression, it is generally important to standardize continuous predictors before training a model:

$$z_j = \frac{x_j - \mu_j}{\sigma_j}$$

Thank you!